

NFL Play Caller

Machine learning model to predict run vs pass plays in NFL games based on pre-snap situational features.
Built using logistic regression on 2021-2023 play-by-play data.

Project Structure

```
nflplaycaller/
    └── src/nfl_run_pass/           # Core ML library
        ├── __init__.py
        ├── config.py
        ├── data_loading.py
        ├── preprocessing.py
        ├── features.py
        ├── models.py
        ├── evaluation.py
        ├── pipeline.py
        └── tuning.py

    └── api/                         # FastAPI backend
        └── main.py                  # REST API for predictions

    └── streamlit/                  # Frontend applications
        ├── st_app.py
        ├── streamlit_btm.py
        └── streamlit_app.py

    └── models/                      # Saved artifacts
        ├── log_reg_model.pkl
        ├── scaler.pkl
        └── feature_cols.json

    └── data/                        # Data files (gitignored)
        └── raw/
            └── pbp_2021_2023.csv  # NFL play-by-play dataset

    └── tests/                      # Unit and integration tests (TO BE CREATED)
        ├── test_data_loading.py
        ├── test_features.py
        ├── test_models.py
        ├── test_pipeline.py
        └── test_api.py

    └── run_pipeline.py
    └── sample_plays.csv
    └── requirements.txt
    └── pyproject.toml
    └── README.md                  # This file
```

Module Overview

Core Library (`src/nfl_run_pass/`)

The library follows a modular pipeline design:

1. `config.py` - Central configuration

- Data paths and filtering rules
- Feature definitions
- Model hyperparameters
- Train/test split settings

2. `data_loading.py` - Data ingestion

- `load_raw_data()` - Read CSV
- `filter_season()` - Filter to 2023 season
- `filter_run_pass_plays()` - Keep only run/pass plays
- `add_is_pass_target()` - Create binary target (0=run, 1=pass)

3. `features.py` - Feature engineering

- `add_engineered_features()` - Create pre-snap features:
 - Field position (red zone, goal-to-go, yardline)
 - Down/distance buckets (short/medium/long)
 - Formation (shotgun, no huddle)
 - Score state (trailing, tied, leading)
 - Time context (quarter, late half)
- `build_feature_matrix()` - Construct (X, y) for modeling

4. `preprocessing.py` - Data preparation

- `handle_missing_values()` - Imputation
- `split_train_test()` - Stratified train/test split
- `scale_features()` - StandardScaler normalization

5. `models.py` - Model training

- `create_base_log_reg_model()` - Logistic regression setup
- `train_log_reg_model()` - Fit with optional GridSearchCV

6. `evaluation.py` - Model assessment

- Accuracy, precision, recall, F1, ROC-AUC
- Confusion matrix
- Train vs test metrics

7. `pipeline.py` - Orchestration

- `run_training_pipeline()` - End-to-end workflow:
 1. Load data
 2. Engineer features

3. Split and scale
4. Train model
5. Evaluate
6. Save artifacts
 - o `save_artifacts()` / `load_artifacts()` - Persistence
8. `tuning.py` - Hyperparameter optimization
 - o `run_log_reg_tuning()` - GridSearchCV for C and class_weight
 - o `compare_models()` - Baseline vs tuned comparison

API (`api/`)

FastAPI REST endpoint for real-time predictions:

- **POST `/predict`** - Takes game situation JSON, returns prediction + probabilities
- Loads saved model/scaler on startup
- CORS-enabled for frontend integration

Frontend (`streamlit/`)

Three Streamlit applications:

1. `st_app.py` - Main interactive predictor with visual field
2. `steamlit_bttn.py` - "Beat the Model" game with real plays + GIFs
3. `streamlit_app.py` - Alternative layout

Setup Instructions

1. Install Dependencies

Using pip:

```
pip install -r requirements.txt
```

Using uv (recommended):

```
uv pip install -e .
```

2. Prepare Data

Place your play-by-play CSV at:

```
data/raw/pbp_2021_2023.csv
```

The data should contain nflfastR-style columns including:

- `season`, `week`, `game_id`, `play_id`
- `play_type` (run/pass/etc)
- `down`, `ydstogo`, `yardline_100`
- `shotgun`, `no_huddle`, `qtr`
- `posteam_score`, `defteam_score`, `score_differential`
- `half_seconds_remaining`, `game_seconds_remaining`
- `posteam`, `home_team`, `away_team`

3. Train the Model

```
python run_pipeline.py
```

This will:

- Load and filter 2023 season data
- Engineer pre-snap features
- Split train/test (80/20, stratified)
- Scale features with StandardScaler
- Train logistic regression with GridSearchCV
- Save model artifacts to `models/`

4. Run Applications

Streamlit UI:

```
streamlit run streamlit/st_app.py
```

FastAPI Backend:

```
cd api  
uvicorn main:app --reload
```

Model Details

Features (18 total)

Base numeric (4):

- `down` (1-4)
- `ydstogo` (yards to first down)
- `yardline_100` (distance to opponent end zone)
- `game_seconds_remaining` (time left in half)

Engineered binary (14):

- Field position: `is_red_zone`, `is_goal_to_go`
- Distance buckets: `short_ydstogo` (≤ 3), `medium_ydstogo` (4-7), `long_ydstogo` (≥ 8)
- Formation: `shotgun`, `no_huddle`
- Score state: `is_trailing`, `is_tied`, `is_leading`, `score_differential`
- Time: `is_fourth_qtr`, `late_half` (<2 min)
- Home/away: `is_home_offense`

Model Architecture

- **Algorithm:** Logistic Regression
- **Preprocessing:** StandardScaler on all features
- **Hyperparameter Tuning:** GridSearchCV
 - `C`: [0.01, 0.1, 1.0, 10.0]
 - `class_weight`: [None, "balanced"]
 - Scoring: F1
 - CV folds: 3

Performance

Add your metrics here after training

```
Train accuracy: X.XX%
Test accuracy: X.XX%
Test F1: X.XX
Test ROC-AUC: X.XX
```

Testing (TODO - Priority #1)

We need to create comprehensive tests in `tests/`:

Unit Tests Needed

1. `test_data_loading.py`
 - Test CSV loading with valid/invalid paths
 - Test season filtering
 - Test run/pass filtering
 - Test target creation
2. `test_features.py`
 - Test each engineered feature calculation
 - Test missing value handling
 - Test feature matrix construction
 - Test feature alignment with saved `feature_cols.json`
3. `test_preprocessing.py`
 - Test train/test split stratification

- Test scaling (mean=0, std=1)
- Test missing value imputation

4. `test_models.py`

- Test model creation
- Test training with/without GridSearchCV
- Test prediction shape and probabilities

5. `test_pipeline.py`

- Test end-to-end pipeline
- Test artifact saving/loading
- Test reproducibility with random_state

6. `test_evaluation.py`

- Test metric calculations
- Test edge cases (all one class, etc)

7. `test_api.py`

- Test /predict endpoint
- Test input validation
- Test feature reconstruction matches training

Integration Tests

- Full pipeline on `sample_plays.csv`
- API + frontend integration
- Model versioning and reproducibility

Testing Framework

Use pytest:

```
pip install pytest pytest-cov
pytest tests/ -v
pytest tests/ --cov=src/nfl_run_pass --cov-report=html
```

Configuration

All settings are in `src/nfl_run_pass/config.py`:

```
from nfl_run_pass.config import CONFIG

# Access configuration
CONFIG.data.season      # 2023
```

```
CONFIG.train_test.test_size # 0.2
CONFIG.log_reg.param_grid # GridSearchCV params
```

To modify:

```
CONFIG.data.season = 2022
CONFIG.train_test.test_size = 0.25
```

Development Workflow

For New Features

1. Update `config.py` with any new settings
2. Modify relevant module (e.g., `features.py`)
3. Write tests for new functionality
4. Run pipeline to retrain model
5. Update API/frontend if needed

For Bug Fixes

1. Write a failing test that reproduces the bug
2. Fix the bug
3. Verify test passes
4. Check integration tests still pass

Before Committing

```
# Run all tests
pytest tests/ -v

# Check code formatting
black src/ tests/
isort src/ tests/

# Run type checking (optional)
mypy src/
```

API Usage

Request Format

```
{
  "down": 3,
  "ydstogo": 7,
  "yardline_100": 35,
```

```
"offense_score": 14,  
"defense_score": 17,  
"qtr": 4,  
"seconds_remaining_half": 180,  
"shotgun": true,  
"no_huddle": false,  
"is_home_offense": true  
}
```

Response Format

```
{  
  "prediction": "PASS",  
  "prob_pass": 0.73,  
  "prob_run": 0.27  
}
```

Example with curl

```
curl -X POST "http://localhost:8000/predict" \  
-H "Content-Type: application/json" \  
-d '{  
  "down": 3,  
  "ydstogo": 7,  
  "yardline_100": 35,  
  "offense_score": 14,  
  "defense_score": 17,  
  "qtr": 4,  
  "seconds_remaining_half": 180,  
  "shotgun": true,  
  "no_huddle": false,  
  "is_home_offense": true  
}'
```

Contributing

Getting Started

1. Clone the repository
2. Create a virtual environment: `python -m venv venv`
3. Install dependencies: `pip install -r requirements.txt`
4. Install package in editable mode: `pip install -e .`
5. Run tests: `pytest tests/`

Branching Strategy

- `main` - Production-ready code

- `develop` - Integration branch
- `feature/feature-name` - New features
- `fix/bug-description` - Bug fixes
- `test/test-description` - Test additions

Pull Request Process

1. Create feature branch from `develop`
2. Write code + tests
3. Ensure all tests pass
4. Update documentation if needed
5. Submit PR to `develop`
6. Get code review from at least one teammate
7. Merge after approval

Team Collaboration

Code Review Checklist

- Tests written for new functionality
- All tests pass
- Code follows project style (see existing modules)
- Documentation updated (docstrings, README)
- No sensitive data or credentials in code
- Configuration changes documented
- Breaking changes clearly communicated

Communication

- Use GitHub Issues for bugs and feature requests
- Use GitHub Projects for sprint planning
- Document major decisions in `/docs/decisions/`
- Update this README as project evolves

Resources

- **nflfastR data guide:** https://www.nflfastr.com/articles/beginners_guide.html
- **scikit-learn docs:** <https://scikit-learn.org/stable/>
- **FastAPI tutorial:** <https://fastapi.tiangolo.com/tutorial/>
- **Streamlit docs:** <https://docs.streamlit.io/>

License

[Add your license here]

Authors

[Add team member names and roles]

Questions? Open an issue or reach out to the team on [your communication channel].